



# Learning in a Small World

B. Ravindran

Reconfigurable and Intelligent Systems Engineering Group

Department of Computer Science and Engineering

Indian Institute of Technology Madras



# Overview

---

## Reinforcement Learning

- Stochastic optimization to solve sequential decision problems
- Hierarchical Decomposition to speed up solution
  - Reuse solutions
  - Reduce number of decisions needed
- What are the appropriate “sub-problems” to learn?
  - Past work uses ideas from social network analysis, e.g. betweenness

## Small World Networks

- Interested in distributed navigation property
- “Long range links” inversely proportional to some power of span
  - Depends on order of lattice

## This Work

- Solutions to sub-problems are equivalent to long-range links
- Create small-world decision problems
- Efficient Learning



# What is Reinforcement Learning?


- Learning about stimuli and actions based on rewards and punishments alone.
- No detailed supervision available
- Trial-and-error learning
- Delayed rewards
- Sequence of actions required to obtain reward
- Associative learning required
  - Need to associate actions to states
- Learn about policies not just actions
- Typically in a stochastic world





# Running a Maze


Dayan and Abbott





# Running a Stochastic Maze

Dayan and Abbott





# The Agent Learns a Policy

---

**Policy** at step  $t$ ,  $\pi_t$  :

a mapping from states to action probabilities

$\pi_t(s, a)$  = probability that  $a_t = a$  when  $s_t = s$

- Reinforcement learning methods specify how the agent changes its policy as a result of experience.
- Roughly, the agent's goal is to get as much reward as it can over the long run.



# The Markov Property

- “the state” at step  $t$ , means whatever information is available to the agent at step  $t$  about its environment.
- The state can include immediate “sensations,” highly processed sensations, and structures built up over time from sequences of sensations.
- Ideally, a state should summarize past sensations so as to retain all “essential” information, i.e., it should have the **Markov Property**:

$$\Pr\{s_{t+1} = s', r_{t+1} = r \mid s_t, a_t, r_t, s_{t-1}, a_{t-1}, \dots, r_1, s_0, a_0\} = \Pr\{s_{t+1} = s', r_{t+1} = r \mid s_t, a_t\}$$

for all  $s', r$ , and histories  $s_t, a_t, r_t, s_{t-1}, a_{t-1}, \dots, r_1, s_0, a_0$ .



# Markov Decision Processes

- If a reinforcement learning task has the Markov Property, it is basically a **Markov Decision Process (MDP)**.
- If state and action sets are finite, it is a **finite MDP**.
- To define a finite MDP, you need to give:
  - **state and action sets**
  - one-step “dynamics” defined by **transition probabilities**:

$$P_{ss'}^a = \Pr\{s_{t+1} = s' \mid s_t = s, a_t = a\} \quad \text{for all } s, s' \in S, a \in A(s).$$

- **reward expectations**:

$$R_{ss'}^a = E\{r_{t+1} \mid s_t = s, a_t = a, s_{t+1} = s'\} \quad \text{for all } s, s' \in S, a \in A(s).$$






# Markov Decision Processes

---


- MDP,  $M$ , is the tuple:  $M = \langle S, A, \Psi, P, R \rangle$ 
  - $S$  : set of states.
  - $A$  : set of actions.
  - $\Psi \subseteq S \times A$  : set of admissible state-action pairs.
  - $P : \Psi \times S \rightarrow [0,1]$  : probability of transition.
  - $R : \Psi \rightarrow \mathbb{R}$  : expected reward.
- Policy  $\pi : S \rightarrow A$  (can be stochastic)
- Maximize total expected reward



# Example



$$M = \langle S, A, \Psi, P, R \rangle$$





# Value Functions

- The **value of a state** (action) is the expected long-term reward starting from that state (action); depends on the agent's policy:

**State - value function for policy  $\pi$  :**

$$V^\pi(s) = E_\pi \{R_t \mid s_t = s\} = E_\pi \left\{ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s \right\}$$

**Action - value function for policy  $\pi$  :**

$$Q^\pi(s,a) = E_\pi \{R_t \mid s_t = s, a_t = a\} = E_\pi \left\{ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s, a_t = a \right\}$$



# Optimal Value Functions


---

- For finite MDPs, policies can be partially ordered:  
 $\pi \geq \pi'$  if and only if  $V^\pi(s) \geq V^{\pi'}(s)$  for all  $s \in S$
- There is always at least one (and possibly many) policies that is better than or equal to all the others. This is an optimal policy. We denote them all  $\pi^*$ .
- Optimal policies share the same optimal state-value function:


$$V^*(s) = \max_{\pi} V^\pi(s) \quad \text{for all } s \in S$$



# Example



$$M = \langle S, A, \Psi, P, R \rangle$$





# Bellman Optimality Equation for $V^*$

$$\begin{aligned} V^*(s) &= \max_{a \in A(s)} E\{r_{t+1} + \gamma V^*(s_{t+1}) \mid s_t = s, a_t = a\} \\ &= \max_{a \in A(s)} \sum_{s'} P_{ss'}^a [R_{ss'}^a + \gamma V^*(s')] \end{aligned}$$

$V^*$  is the unique solution of this system of nonlinear equations.

Similar results applicable to  $Q^*$




# Solution Methods

---

- Temporal Difference Methods
  - TD( $\lambda$ )
  - Q-learning
  - SARSA
  - Actor-Critic
- Policy Search
  - Policy Gradient Methods
  - Evolutionary algorithms
- Stochastic Dynamic Programming



# Sub-goals based decomposition



Features:

rooms = {0, 1, 2, 3, 4, 5}

x = {0, ..., 9}

y = {0, ..., 19}

binary:  $have_i$   $i = 1, \dots, 5$

- Task is to collect all objects in the world
- 5 sub-tasks – one for each room.





# Hierarchical Reinforcement Learning

Options (Sutton, Precup, & Singh, 1999):


An option is a triple  $o = \langle I, \pi_o, \beta \rangle$

- $I \subseteq S$  is the set of states in which  $o$  may be started
- $\pi_o: \Psi \rightarrow [0,1]$  is the (stochastic) policy followed during  $o$
- $\beta: S \rightarrow [0,1]$  is the probability of terminating in each state



# Generalising over Tasks


- Each task has a different reward structure in the state space





# Generalising over Tasks

- Each task has a different reward structure in the state space
- Options provide a model for subtasks.





# Discovering skills

---

- Bottlenecks [McGovern & Barto, Stolle & Precup, etc.]
- Graph partitions [Mannor et al., Mathew et al.]
- Betweenness [Simsek & Barto]
- Frequency of changes [Jonsson & Barto, Hengst]
- Bisimulation metrics [Castro & Precup]
- **Our Hypothesis:** Choose options so that navigating greedily using the Q function would be more effective
  - Turn the MDP into a “small” world



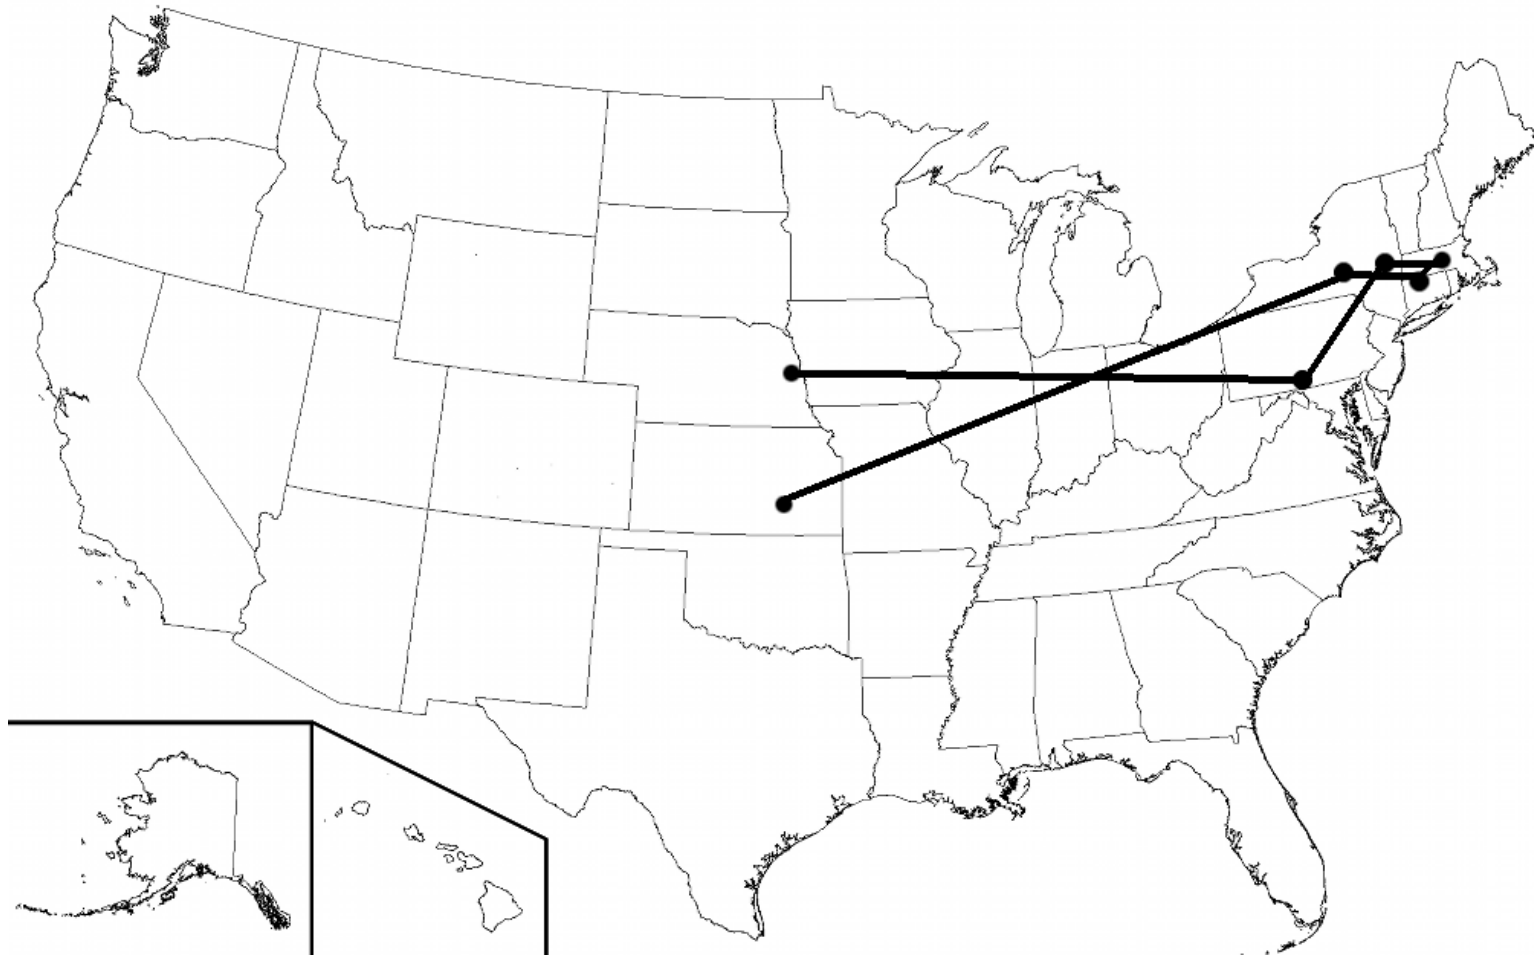
# Small Worlds

---

- A network model where every node is reachable in  $O(\log(n))$  hops from a node.
- More importantly, the  $O(\log(n))$  path can be discovered using a decentralised algorithm.
- Stanley Milgram's Experiment



# Small Worlds





# Small Worlds


---

- A network model where every node is reachable in  $O(\log(n))$  hops from a node.
- More importantly, the  $O(\log(n))$  path can be discovered using a decentralised algorithm.
- Stanley Milgram's Experiment
- Kleinberg's key insight: Power-law distributed neighbours



# Small Worlds

- A  $r$ -dimensional lattice graph
- Edges distributed inversely proportional to distance

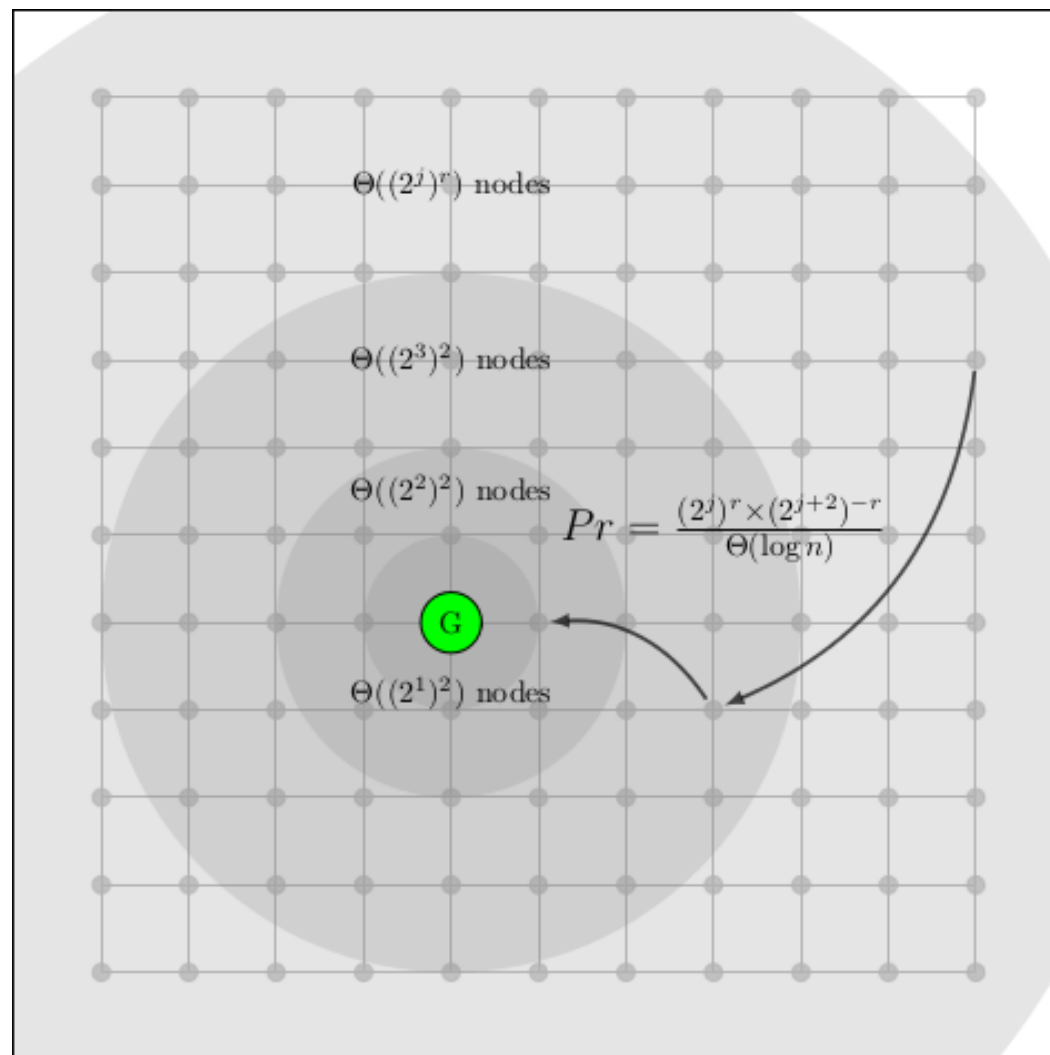






# Small Worlds

- A  $r$ -dimensional lattice graph
- Edges distributed inversely proportional to distance
- A greedy agent will move from one neighbourhood to another in  $\log(n)$  time






# Small Worlds in RL

- Construct “path options” that take an agent from state  $s$  to  $s'$ .
- $s'$  is chosen according to the power-law.
- Which distance based?
- Value and state-space distance are related

$$k_1 \|u - v\| - c_1 \leq \|u - v\|_v \leq k_2 \|u - v\|$$





# Small Worlds in RL

---

- Many options required; how do we effectively learn them?
  - *Note: We add only one extra option per state*
- Key Insight: The important point is to move to an exponentially smaller neighbourhood of target.
  - Use cheap, possibly inaccurate options
- Algorithm:
  - Train an agent on  $T$  different tasks.
  - For each task, save path options using two states distributed according to the power-law, and following the policy gradient.
- Does not require complete knowledge of the MDP, nor does it need to build a model



# Algorithm

---

## Algorithm 2 QOptions: Options from a $Q$ -Value Function

---

**Require:**  $Q, r, n$

- 1:  $O \leftarrow \emptyset$
  - 2:  $\pi \leftarrow$  greedy policy from  $Q$
  - 3: **for all**  $s$  in  $S$  **do**
  - 4:   Choose an  $s'$  according to  $P_r$
  - 5:   **if**  $Q(s', \pi(s')) > Q(s, \pi(s))$  **then**
  - 6:      $O \leftarrow O \cup \langle \{s\}, \pi, \{s'\} \cup \{t \mid Q(s', \pi(s')) < Q(t, \pi(t)) \} \rangle$
  - 7:   **end if**
  - 8: **end for**  $s$  in  $S$
  - 9: **return** A random subset of  $n$  options from  $O$
-



# Experiments

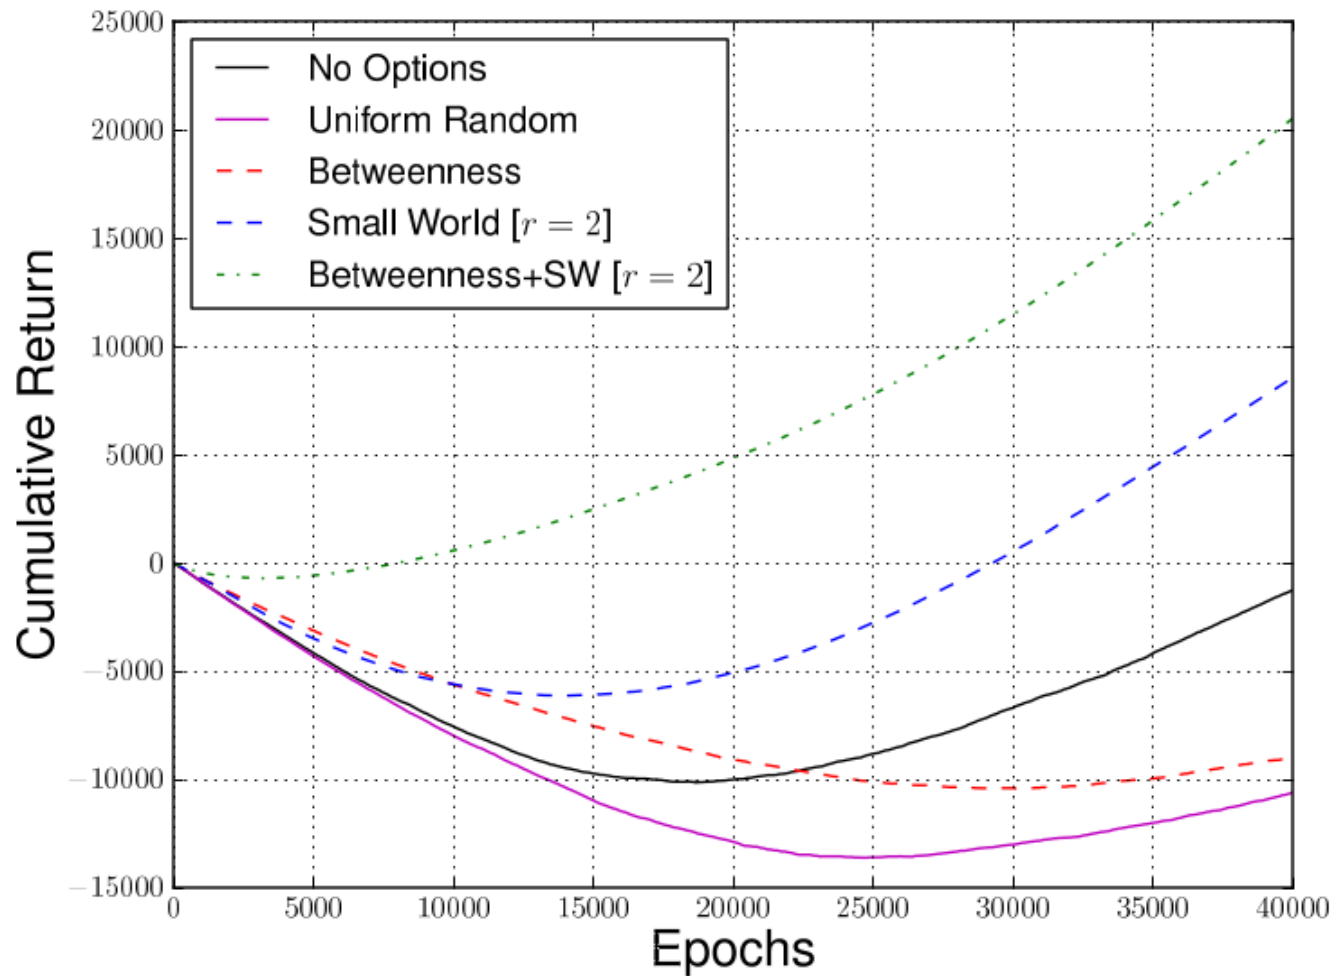
---

1. No options
  2. 200 options generated uniformly randomly
  3. Betweenness based options
  4. 200 small world options
  5. Betweenness + Small World
-



# Experimental Results

Rooms: Cumulative Return with 200 options





# Experimental Results

	Arbitrary Navigation	Rooms	Taxi
None	-31.82	-1.27	-16.90
Random	-31.23	-10.76	-18.83
Betweenness	-18.28	-8.94	<b>80.48</b>
Small World	<b>-14.24 [r = 4]</b>	<b>8.54 [r = 2]</b>	0.66 [r = 0.75]

R				G
Y			B	



# Learning on a budget

---


- Differing amounts of effort required to learn options
  - Fewer betweenness options than small world options
  - Equalize budget for both settings
    - $10^6$  epochs
    - $10^7$  epochs
    - Optimal solutions
  - Recall: It is sufficient if the small world option “hops” to closer region
-





# Experimental Results

## Rooms: Options Learnt on a Budget





# Conclusions

---


- A new option generation scheme with theoretically motivated generalisation properties.
- An efficient algorithm to learn “small world” options solely from experience on a few tasks.
- Competitive performance on standard domains



# Future Work

- Extending small world options to continuous domains
- Experiments on more diverse domains
- What factors affect the power-law exponent  $r$ ?
- Adding options online
- Exploring implications on learning bounds

Rooms:  $r$  vs. Cumulative Return





# Questions?

---

- <http://www.cse.iitm.ac.in/~ravi>
-